

LLM segmentation needs a large context buffer; Hi-OnTop runs at 0-latency
AMI meetings (n=12). LLM gains only emerge with near-full context; at streaming buffers it ties cheap 0-latency methods.

